

Introduction to the t-test: data exploration

Nathan Brouwer / brouwern@gmail.com / @lobrowr

September 29, 2016

Introduction

In this lab exercise we will use data from Chapter 12 of Whitlock & Shulter's *Analysis of Biological Data*. The data are from an observational study (Levin et al 2002) where the survival of native Chinook Salmon was monitored in 6 streams in the Columbia River Basin in Idaho where the eastern species brook trout had been introduced for sport fishing. Brook trout were hypothesized to reduce survival rates of the native Chinook salmon. Six sites where brook trout were absent were used as controls. In each stream, juvenile Chinook were captured, PIT tagged, and released. PIT tags are similar to the microchips put in pets to ID them. As the fish migrated downstream, their PIT tags were detected by receivers. The number of PIT tags detected divided by the total number released indicates the survival rate. The original study used several years of data; Whitlock and Shulter present just one, which we will examine in detail.

References

Levin et al. 2002. Non-indigenous brook trout and the demise of Pacific salmon: a forgotten threat? Proceedings of the Royal Society B. <http://rspb.royalsocietypublishing.org/content/269/1501/1663.short>

Preliminaries

Load packages

```
# For plots
library(ggplot2)

# Great ggplot helper functions
library(ggpubr)

# data manipulation
library(dplyr)
```

Set the number of digits to display

R's default is to display numbers to a high level of precision: 7 digits. We can see this using this command

```
options()$digits
```

```
## [1] 7
```

We can set R to automatically round things off like this

```
options(digits = 4)
```

Working directory

Set working the directory if needed. We'll be loading data "by hand" so this probably isn't necessary

Load Brook Trout Data

These data are from Example 12.4 “So long; thanks to all the fish” in Whitlock and Shulter 2nd ed.

Load data by hand

We'll 1st load three pieces of info into vectors then put them together into a dataframe with 3 columns

- 1) Make A vector to hold the status of the stream - whether brook trout are present or absent

```
brook.trout.PRES.ABS <- c("present", "present", "present", "present", "present", "present",  
                          "absent", "absent", "absent", "absent", "absent", "absent")
```

Advanced technique

This is **OPTIONAL**: A nice trick is this, using the rep() command. Note the c() that surrounds BOTH of the rep() commands.

```
brook.trout.PRES.ABS <- c(rep("present", 6), rep("absent", 6))
```

- 2) The number of salmon released

```
salmon.released <- c(820, 960, 700, 545, 769, 1001, 467, 959, 1029, 27, 998, 936)
```

- 3) The number of salmon that survived

```
salmon.surv <- c(166, 136, 153, 103, 173, 188, 180, 178, 326, 7, 120, 135)
```

- 4) Make a **dataframe** from the 3 vectors

```
salmon <- data.frame(brook.trout.PRES.ABS,  
                    salmon.released,  
                    salmon.surv)
```

The resulting dataframe has 12 rows and 3 columns

```
dim(salmon)
```

```
## [1] 12  3
```

OPTIONAL: Load via .csv file

If the data are in a .csv file in the working directory they could be loaded as. This does NOT need to be done if you have already loaded it by hand as above.

```
salmon <- read.csv("Lab7_data_brook_trout.csv")
```

Examine the brook trout data

Each line of data is a separate stream that either has introduced brook trout present or not present.

```
# size of dataframe  
dim(salmon)  
  
#top of data  
head(salmon)  
  
#bottom of data  
tail(salmon)
```

```
#summary of data
summary(salmon)
```

Data setup

- The **response variable** in this study is the **percent (%)** of chinook salmon that survived over the course of the study.
- We therefore need to calculate this value using R's basic math function for division.

Calculate survival rates

1st, calculate the fraction of fish that survived in each stream

- Note the use of the "\$" operator to access the columns
- We access the survival column using "salmon\$salmon.surv" and divide this by the number released in "salmon\$salmon.released"

```
percent.surv <- salmon$salmon.surv / salmon$salmon.released
```

2nd, check to make sure that the data make sense.

- The data are a percentage, so each datum should be between 0 and 1.
- We can look at the data using `summary()`
- Note that we can use `summary` on a single vector (`percent.surv`) or on a dataframe (`salmon`)

```
summary(percent.surv)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    0.120   0.175   0.196   0.215   0.234   0.385
```

We could also use `range()`, `min()`, and `max()` to get this info. We could also make a histogram using `hist()`.

3rd, add the percentages to the original dataframe

- We'll call the new column "percent.surv" also.
- We make a new column by typing "salmon\$percent.surv"
- "<- percent.surv" adds the "percent.surv" vector to this new column in the dataframe

```
salmon$percent.surv <- percent.surv
```

Note: this isn't actually a percent, but a fraction. To get a percent we'd multiply it by 100.

4th, check the calculated data

```
#Look at the whole dataframe
salmon
```

```
##      brook.trout.PRES.ABS salmon.released salmon.surv percent.surv
## 1                present             820          166      0.2024
## 2                present             960          136      0.1417
## 3                present             700          153      0.2186
## 4                present             545          103      0.1890
```

```
## 5          present          769          173          0.2250
## 6          present         1001          188          0.1878
## 7          absent          467          180          0.3854
## 8          absent          959          178          0.1856
## 9          absent         1029          326          0.3168
## 10         absent           27           7          0.2593
## 11         absent          998          120          0.1202
## 12         absent          936          135          0.1442
```

```
#summary() of whole dataframe
summary(salmon)
```

```
## brook.trout.PRES.ABS salmon.released salmon.surv percent.surv
## absent :6           Min.   : 27    Min.   : 7    Min.   :0.120
## present:6           1st Qu.: 661    1st Qu.:131   1st Qu.:0.175
##                   Median : 878    Median :160   Median :0.196
##                   Mean    : 768    Mean    :155   Mean    :0.215
##                   3rd Qu.: 970    3rd Qu.:178   3rd Qu.:0.234
##                   Max.    :1029    Max.    :326   Max.    :0.385
```

Data exploration


- We should explore the data using boxplots and histograms.
- The **response variable** is the numeric variable “percent.surv” (again, actually a fracation, sorry...)
- The **predictor variable*** is a categorical variable** the indicates whether introduced brook trout are present or absent.

Boxplots

Basic boxplot

- We can make a boxplot using ggpubr
- Note that in ggpubr, the names of variables occur in quotes “...”

```
ggboxplot(data = salmon,
          y = "percent.surv",
          x = "brook.trout.PRES.ABS")
```




Median survival is a bit lower when brook trout are present, but there is not much difference.

Fancier boxplot


- We can color-code the plot using `fill = "brook.trout.PRES.ABS"`
- Again, note the quotes around "brook.trout.PRES.ABS" in both places it occurs

```
ggboxplot(data = salmon,  
  y = "percent.surv",  
  x = "brook.trout.PRES.ABS",  
  fill = "brook.trout.PRES.ABS")#
```




We can change our axes labels using `xlab =` and `ylab =`

```
ggboxplot(data = salmon,  
  y = "percent.surv",  
  x = "brook.trout.PRES.ABS",  
  fill = "brook.trout.PRES.ABS",  
  xlab = "Brook trout",  
  ylab = "Percent Survival (%)")#
```



- Boxplot indicate medians using a thick line. We can indicate where the mean is using "add = mean"
- The mean shows up as a dot.

```
ggboxplot(data = salmon,  
  y = "percent.surv",  
  x = "brook.trout.PRES.ABS",  
  fill = "brook.trout.PRES.ABS",  
  xlab = "Brook trout",  
  ylab = "Percent Survival (%)",  
  add = "mean")#
```



OPTIONAL: Classic Boxplot


For reference, this is the old-school, pre-ggplot way of doing this. This is **just shown for comparison**.

```
boxplot(percent.surv~brook.trout.PRES.ABS,
        data = salmon)
```

Histograms

This dataset is a bit small so a histogram looks a bit awkward unless you change the defaults. The code below makes a pretty good one. I use “facet.by = ...” to split the graph into 2 panels. Note that I set “bins = 5.” What happens when you remove this?

```
gghistogram(data = salmon,
            x = "percent.surv",
            fill = "brook.trout.PRES.ABS",
            bins = 5,
            facet.by = "brook.trout.PRES.ABS")
```



Calculating mean and standard deviation

We are going to using the dplyr package to extract information from our data.

Reminder: you dataframe

If we just highlight the word “salmon” and run it we get our whole dataframe

```
salmon
```

```
##      brook.trout.PRES.ABS salmon.released salmon.surv percent.surv
## 1                present           820           166      0.2024
## 2                present           960           136      0.1417
## 3                present           700           153      0.2186
## 4                present           545           103      0.1890
## 5                present           769           173      0.2250
## 6                present          1001           188      0.1878
## 7                absent           467           180      0.3854
## 8                absent           959           178      0.1856
## 9                absent          1029           326      0.3168
## 10               absent            27            7      0.2593
## 11               absent           998           120      0.1202
## 12               absent           936           135      0.1442
```

dplyr's summarize function

- dplyr has a function called summarize that can calculate means and other things for us.
- we can tell it to summarize a column from the salmon dataframe using a “pipe”, indicated by “%>%”.
- Within the summarize() command we tell it how to summarize the data, mean(), and what column to summarize, percent.surv

```
salmon %>%
  summarize(mean(percent.surv))
```

```
##      mean(percent.surv)
## 1                0.2147
```

2 summaries at the same time

- We can go further and calculate the mean and the SD.
- Remember the comma between the mean() and sd() line
- The mean is the 1st column, and the SD is in the 2nd column.
- (If you can't see the SD, make your console window larger)

```
salmon %>%
  summarize(mean(percent.surv),
            sd(percent.surv))
```

```
##      mean(percent.surv) sd(percent.surv)
## 1                0.2147          0.07585
```

OPTIONAL We can compare these results to these two other ways of doing things

```
#mean on a single column
mean(salmon$percent.surv)
```

```
## [1] 0.2147
```

```
#summary on a single column
summary(salmon$percent.surv)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
```

```
##    0.120    0.175    0.196    0.215    0.234    0.385
```

dplyr's group_by() function

dplyr has a function called group_by() that can separate out data based on a categorical variable. On its own it doesn't do anything fancy, but in general we can tell it to group things from a dataframe like this. (Don't run this code - its output isn't really useful)

```
salmon %>%  
  group_by(brook.trout.PRES.ABS)
```

Again, Note that the "%>% " is called a "pipe" that tells that group() function to work on the salmon dataframe

Using group_by() with summarize()

- We can "pipe" things together in a series of steps.
- 1st, we declare our dataframe, then we use group_by() to split it apart
- 2nd we use summarize(...) to carry out a summary function
- Note that the pipe command %>% occurs twice

```
salmon %>%  
  group_by(brook.trout.PRES.ABS) %>%  
  summarize(mean(percent.surv))
```

```
## # A tibble: 2 x 2  
##   brook.trout.PRES.ABS `mean(percent.surv)`  
##           <fctr>           <dbl>  
## 1         absent           0.2353  
## 2         present           0.1941
```

We can go further and calculate the mean and the SD for both groups. The mean is the middle column, and the SD is in the far right column. (If you can't see the SD, make your console window larger)

```
salmon %>%  
  group_by(brook.trout.PRES.ABS) %>%  
  summarize(mean(percent.surv),  
            sd(percent.surv))
```

```
## # A tibble: 2 x 3  
##   brook.trout.PRES.ABS `mean(percent.surv)` `sd(percent.surv)`  
##           <fctr>           <dbl>           <dbl>  
## 1         absent           0.2353           0.10369  
## 2         present           0.1941           0.02979
```

We can save our output like by assigning it to an R object like this

```
my.summary <- salmon %>%  
  group_by(brook.trout.PRES.ABS) %>%  
  summarize(mean(percent.surv),  
            sd(percent.surv))
```

And then look at it

```
my.summary
```



```
## # A tibble: 2 x 3
##   brook.trout.PRES.ABS `mean(percent.surv)` `sd(percent.surv)`
##           <fctr>           <dbl>           <dbl>
## 1         absent         0.2353         0.10369
## 2         present         0.1941         0.02979
```

Adding labels

- The output of dply is a bit “verbose”
- We can shorten it by adding our own shorter labels within the summarize() function
- This is similar to how data.frame() works

```
my.summary <- salmon %>%
  group_by(brook.trout.PRES.ABS) %>%
  summarize(mean = mean(percent.surv),
            sd = sd(percent.surv))
```

Getting the standard error (SE)

- We need the SD and the sample size to get the standard error (SE)
- The sample size can be determined using the length() function

Getting the sample size

```
my.summary <- salmon %>%
  group_by(brook.trout.PRES.ABS) %>%
  summarize(mean = mean(percent.surv),
            sd = sd(percent.surv),
            N = length(percent.surv))
```

Calculating the SE

```
my.summary$SE <- my.summary$sd/sqrt(my.summary$N)
```


Plotting the mean and error bars

Boxplot and histograms are good for exploring the overall distribution of the data. We typically want to plot the means with error bars when we are doing hypothesis testing. Note that while the SD tells us about how variable the data are, the standard error tells us about how precise our estimates are, and therefore how confident we can be that two means are different from each other (or whether we can be pretty sure they are similar).

Plotting means

First, we can plot the means like this.


```
ggdotplot(data = my.summary,
          y = "mean",
          x = "brook.trout.PRES.ABS")
```



IGNORE THE ERROR “stat_bindot() using bins = 30. Pick better value with binwidth.”


And give them color by adding “fill =”

```
ggdotplot(data = my.summary,
  y = "mean",
  x = "brook.trout.PRES.ABS",
  fill = "brook.trout.PRES.ABS")
```



We can make things a bit bigger like this Even better than just color is color AND symbol, using shape = ...


```
ggdotplot(data = my.summary,
  y = "mean",
  x = "brook.trout.PRES.ABS",
  fill = "brook.trout.PRES.ABS",
  size = 3)
```



Plotting error bars w/ `geom_errorbar()`

We can combine the `ggpubr` command `"ggdotplot()"` and the regular `ggplot2` command `geom_errorbar()` to plot the means and error bars.

```
ggdotplot(data = my.summary,  
  y = "mean",  
  x = "brook.trout.PRES.ABS",  
  fill = "brook.trout.PRES.ABS",  
  size = 3) +  
  geom_errorbar(aes(ymax = mean + SE,  
    ymin = mean - SE))
```



The really wide bars extension on the error bars are goofy, so we can get ride of them by using `"width = 0"` within `geom_errorbar()`.

```
ggdotplot(data = my.summary,  
  y = "mean",  
  x = "brook.trout.PRES.ABS",  
  fill = "brook.trout.PRES.ABS",  
  size = 3) +  
  geom_errorbar(aes(ymax = mean + SE,  
    ymin = mean - SE),  
    width = 0)
```

- And we can make the error bars wide by using `"size ="` within `geom_errorbar()`

```
ggdotplot(data = my.summary,  
  y = "mean",  
  x = "brook.trout.PRES.ABS",  
  fill = "brook.trout.PRES.ABS",  
  size = 3) +  
  geom_errorbar(aes(ymax = mean + SE,  
    ymin = mean - SE),  
    width = 0,  
    size = 2)
```